

软件工程学院  
SOFTWARE ENGINEERING INSTITUTE

# 人工智能导论

## Lecture 6: Training

# Supervised Learning - Review

---

- Given training data  $\{(x_i, y_i)\}_{i=1}^n$  i.i.d. from distribution  $D$

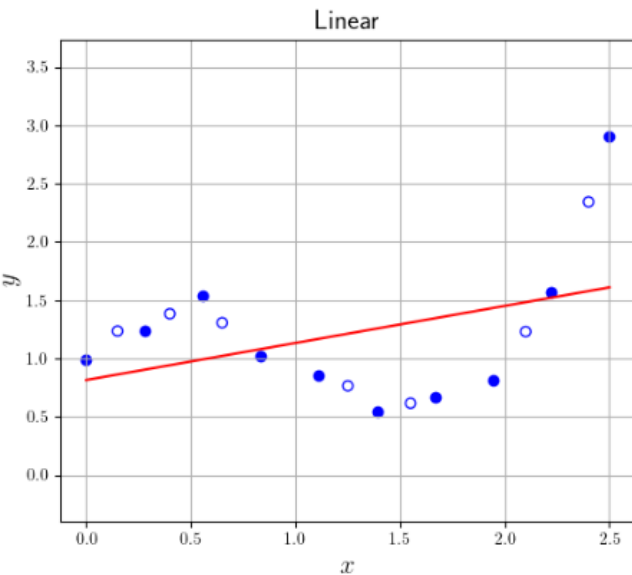


- Find  $y = f(x; \theta)$ 
  - ◆ which works well on **test data** i.i.d. from distribution  $D$
  - ◆  $\theta$  can be trained by minimizing the empirical loss

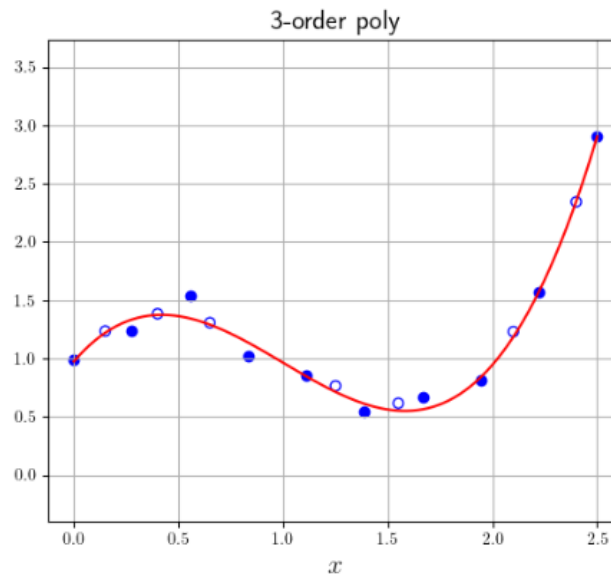
$$\hat{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i; \theta), y_i)$$

# Underfitting & Overfitting

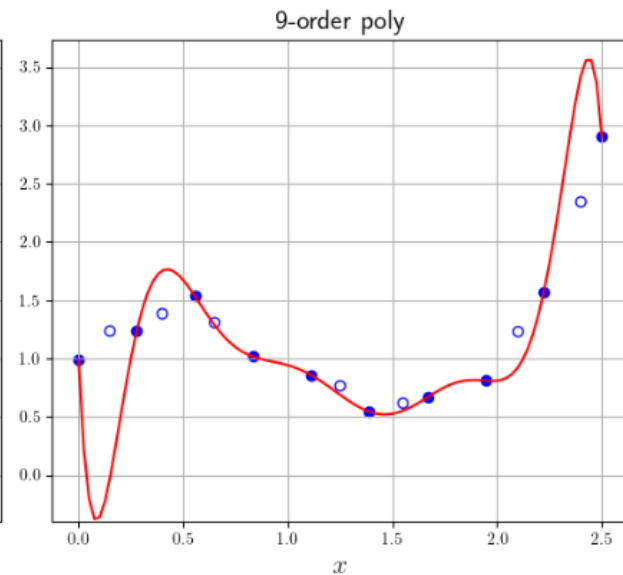
---



underfitting



well fitting



overfitting

# Underfitting & Overfitting

---

- **Underfitting:** the model cannot capture the underlying trend of the data
  - ◆ Large training error
  - ◆ Model is not complex enough
  
- **Overfitting:** the model describes random noise instead of the underlying relationship
  - ◆ Small training error but large test error
  - ◆ Model is too complex

# Outline

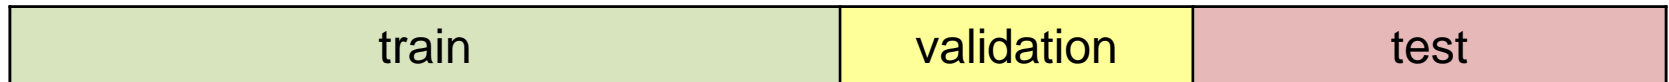
---

- ❑ Model selection
- ❑ Optimization
- ❑ Not work well on training data
- ❑ Not work well on testing data
- ❑ Hyperparameter tuning

# Model selection

---

□ Given  $m$  models  $f_1, f_2 \dots, f_m$ , how to find the best model?



- ◆ Split data into train, validation, and test
- ◆ Choose hyperparameters on the validation data and evaluate on the test data

# Cross validation

---

❑ Split data into folds, try each fold as validation and average the results

❑ Example: 5-fold cross validation

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

# Outline

---

- ❑ Model selection
- ❑ Optimization
- ❑ Not work well on training data
- ❑ Not work well on testing data
- ❑ Hyperparameter tuning



# Optimization

---

- Minimize the empirical loss

$$\min_{\theta} L(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i; \theta), y_i)$$

- Usually  $L(\theta)$  is continuous and differentiable (or subdifferentiable)

# History of optimization

---

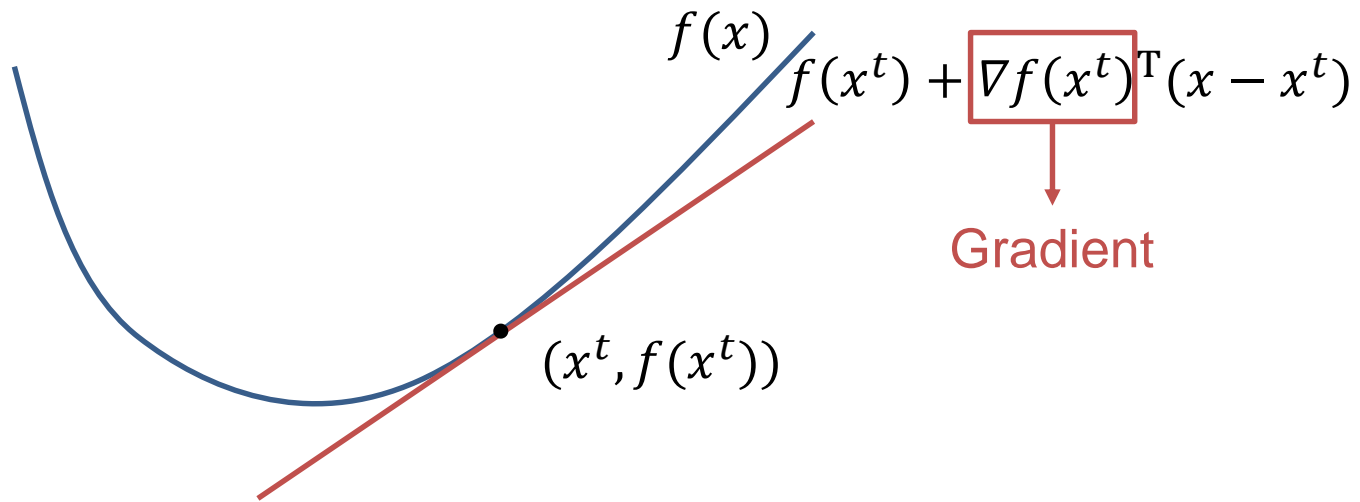
- ◆ 1847: Cauchy proposes gradient descent
- ◆ 1950s: Linear Programs, soon followed by non-linear, Stochastic Gradient Descent (SGD)
- ◆ 1980s: General optimization, convergence theory
- ◆ 2005-2015: Large scale optimization (mostly convex), convergence of SGD
- ◆ 2015-today: Improved understanding of SGD for deep learning

# Gradient descent

---

□ What is the steepest descent direction at  $x^t$ ?

opposite direction of the gradient



# Gradient descent

---

□ Start with an initial point  $x^0$

◆ In each iteration, compute

$$x^{t+1} = x^t - \eta_t \nabla f(x^t)$$

□  $\eta_t$  is the learning rate

# Gradient descent

---

□ Objective function:  $f(\mathbf{x}) = x_1^2 + 2x_2^2$ ,  
initial point:  $\mathbf{x}^{(0)} = (1, 1)^T$ , learning rate  
 $\eta = 0.2$

$$\nabla f(\mathbf{x}) = \begin{bmatrix} 2x_1 \\ 4x_2 \end{bmatrix}$$

$$\mathbf{x}^{(1)} = \mathbf{x}^{(0)} - \eta \nabla f(\mathbf{x}^{(0)}) = \begin{bmatrix} 1 \\ 1 \end{bmatrix} - 0.2 \times \begin{bmatrix} 2 \\ 4 \end{bmatrix} = \begin{bmatrix} 0.6 \\ 0.2 \end{bmatrix}$$

$$\mathbf{x}^{(2)} = \mathbf{x}^{(1)} - \eta \nabla f(\mathbf{x}^{(1)}) = \begin{bmatrix} 0.6 \\ 0.2 \end{bmatrix} - 0.2 \times \begin{bmatrix} 1.2 \\ 0.8 \end{bmatrix} = \begin{bmatrix} 0.36 \\ 0.04 \end{bmatrix}$$

$$\mathbf{x}^{(3)} = \mathbf{x}^{(2)} - \eta \nabla f(\mathbf{x}^{(2)}) = \begin{bmatrix} 0.36 \\ 0.04 \end{bmatrix} - 0.2 \times \begin{bmatrix} 0.72 \\ 0.16 \end{bmatrix} = \begin{bmatrix} 0.216 \\ 0.008 \end{bmatrix}$$

...

# Stochastic gradient descent

---

## □ Gradient descent:

$$\theta^{t+1} = \theta^t - \nabla_{\theta} \left( \frac{1}{n} \sum_{i=1}^n \ell(f(x_i; \theta^t), y_i) \right)$$

## □ Stochastic gradient descent:

◆ Pick an data  $(x_i, y_i)$

$$\theta^{t+1} = \theta^t - \nabla_{\theta} (\ell(f(x_i; \theta^t), y_i))$$

# Mini-batch SGD

---

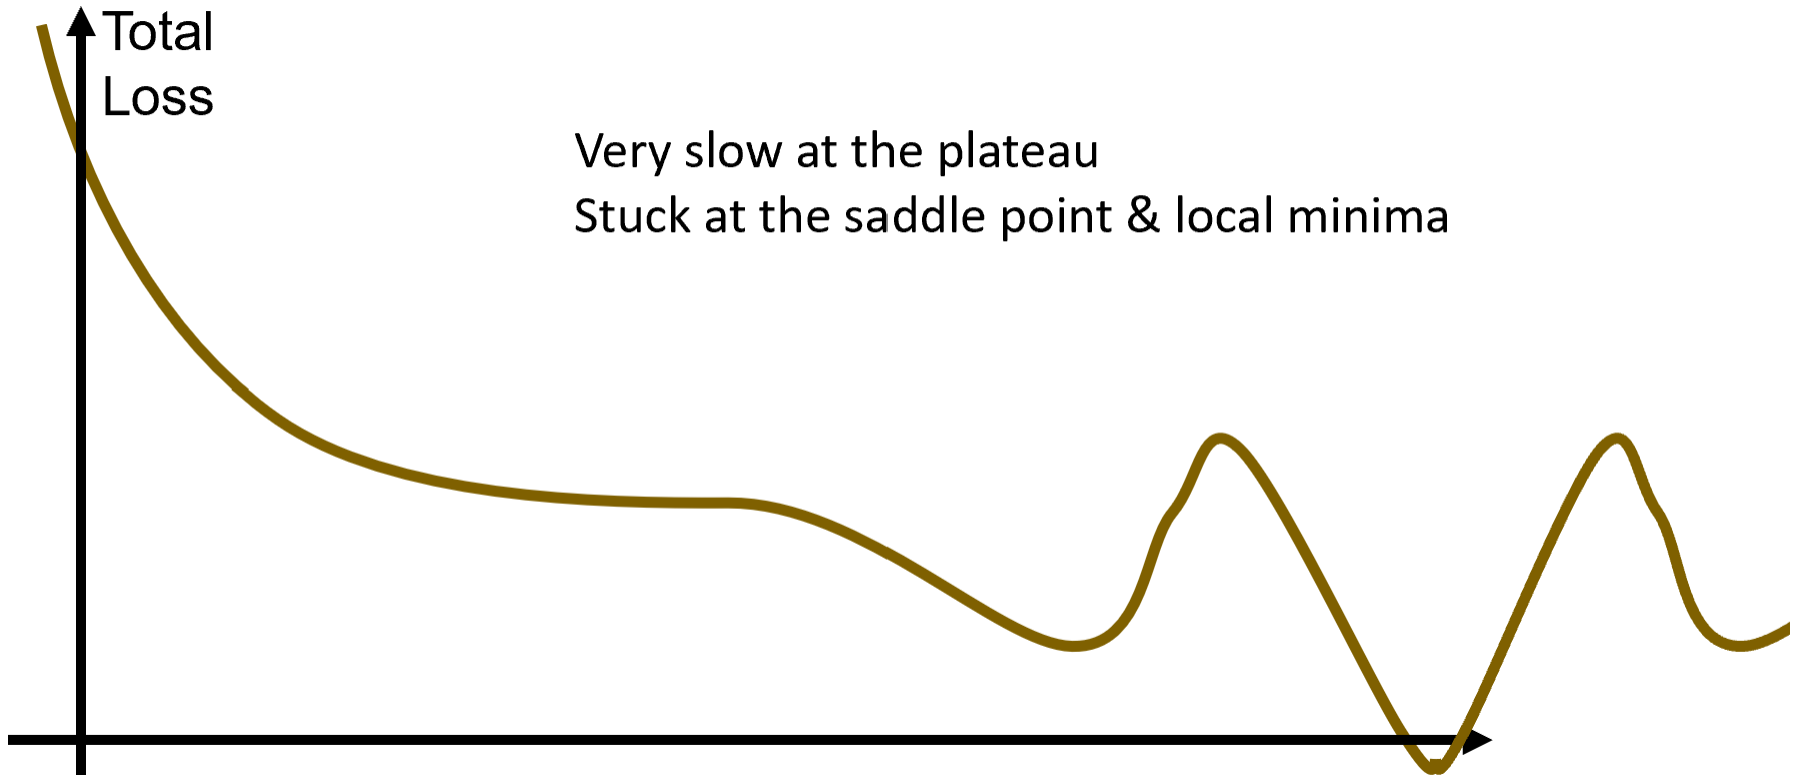
- In each iteration, randomly pick a mini-batch  $S = \{(x_{b_1}, y_{b_1}), \dots, (x_{b_k}, y_{b_k})\}$

$$\theta^{t+1} = \theta^t - \nabla_{\theta} \left( \frac{1}{|S|} \sum_{(x_i, y_i) \in S} \ell(f(x_i; \theta^t), y_i) \right)$$

- What are the parameters?
  - ◆ Learning rate
  - ◆ Batch size
  - ◆ When to stop

# Challenges of nonconvex optimization

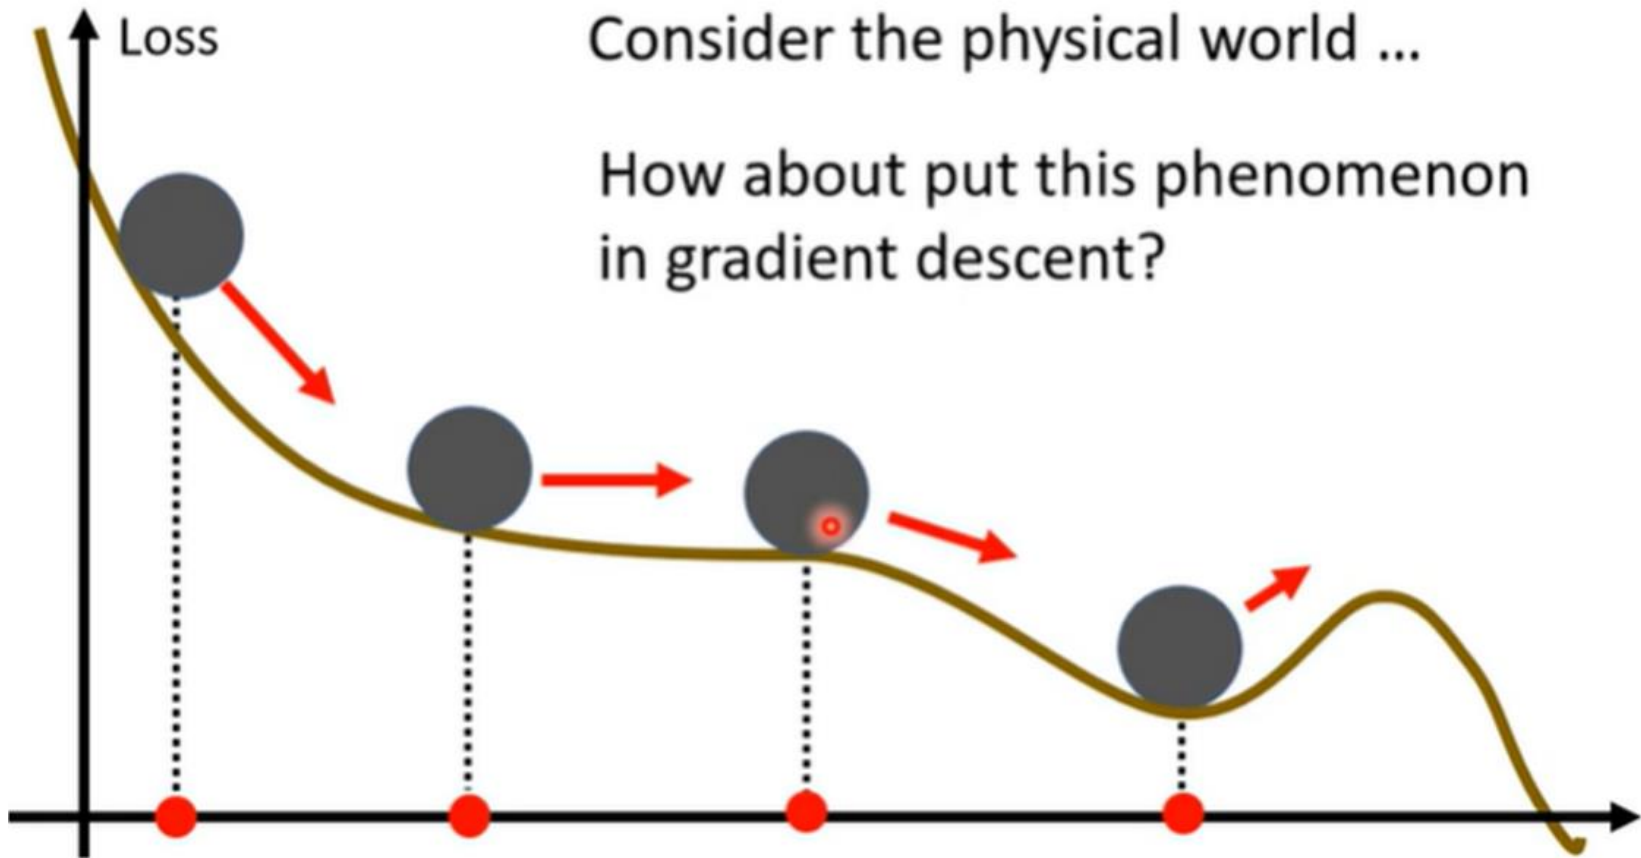
---





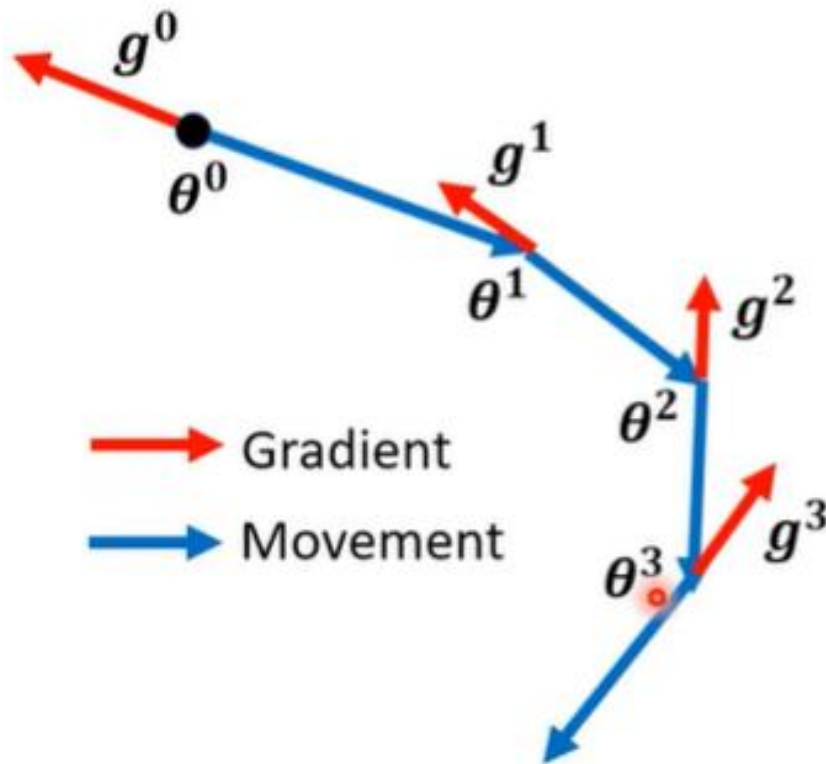
# Momentum

---



# Vanilla Gradient Descent

---



Starting at  $\theta^0$

Compute gradient  $g^0$

Move to  $\theta^1 = \theta^0 - \eta g^0$

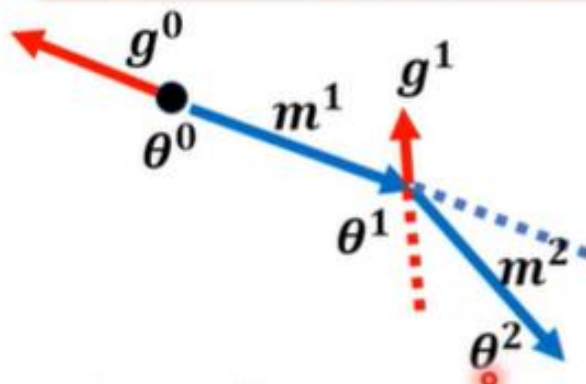
Compute gradient  $g^1$

Move to  $\theta^2 = \theta^1 - \eta g^1$

$\vdots$

# Momentum

Movement: movement of last step minus gradient at present



- Gradient
- Movement
- ..... Movement of the last step

Starting at  $\theta^0$

Movement  $m^0 = 0$

Compute gradient  $g^0$

Movement  $m^1 = \lambda m^0 - \eta g^0$

Move to  $\theta^1 = \theta^0 + m^1$

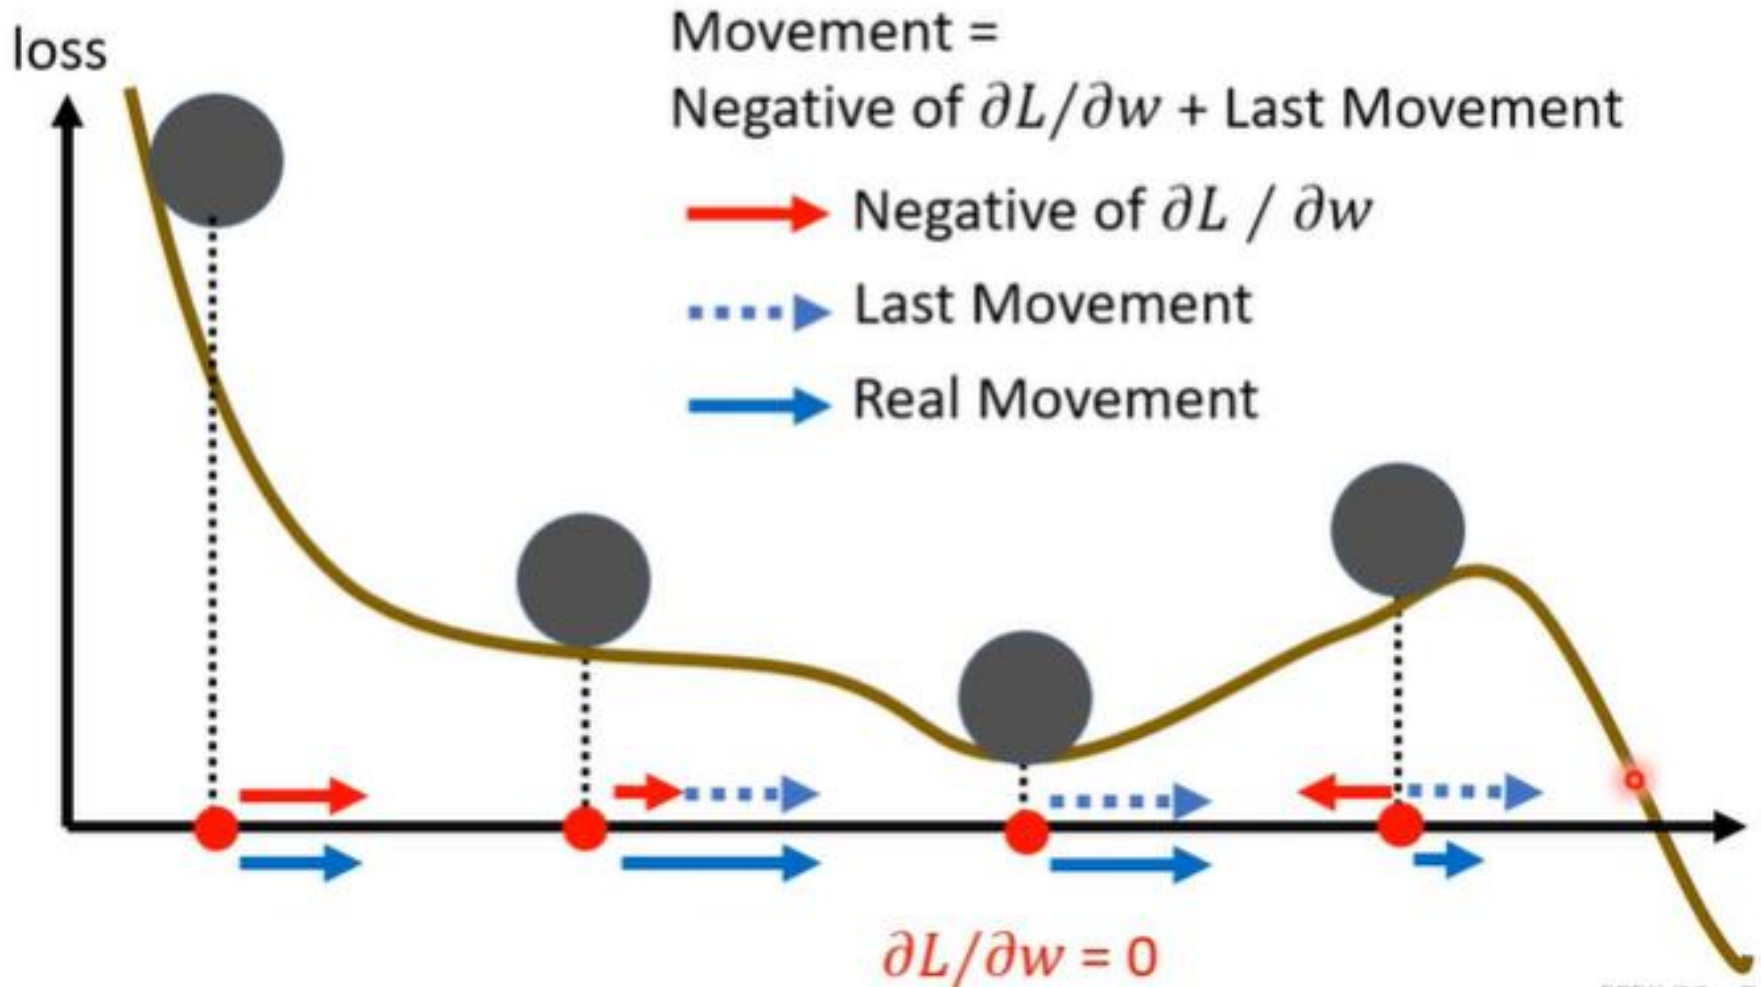
Compute gradient  $g^1$

Movement  $m^2 = \lambda m^1 - \eta g^1$

Move to  $\theta^2 = \theta^1 + m^2$

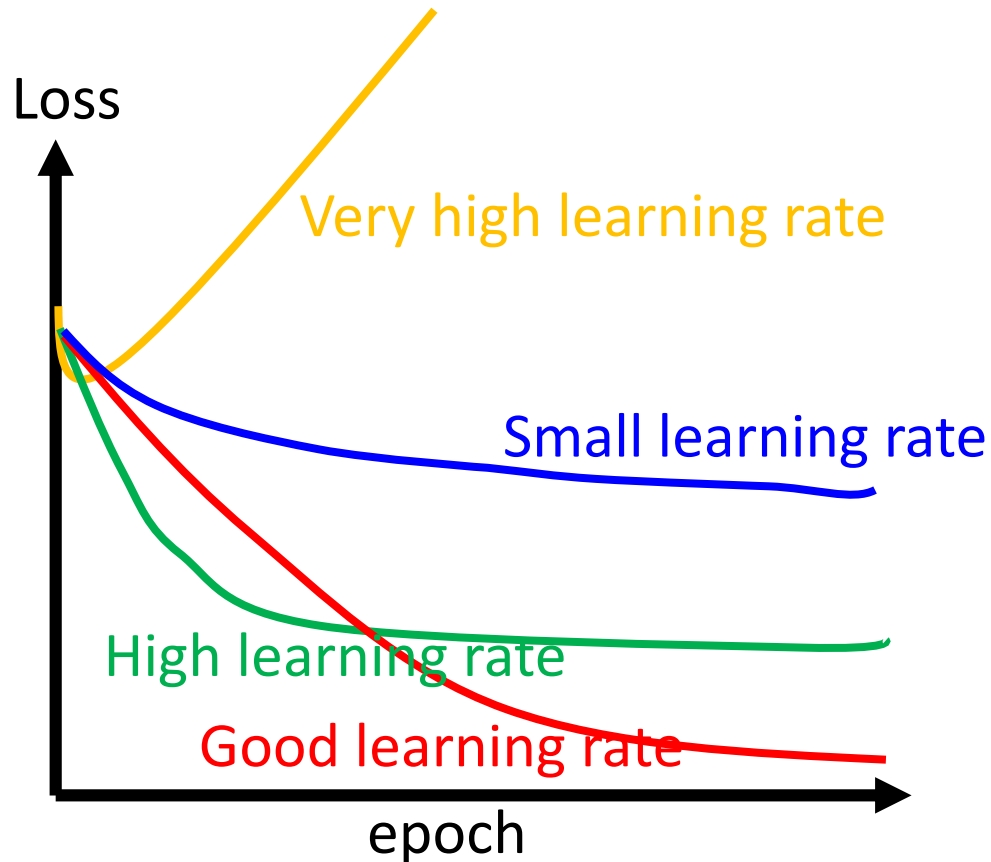
Movement not just based on gradient, but previous movement.

# Momentum



# The effects of different learning rates

---



□ SGD usually require decaying learning rate:

$$\eta_t = \frac{1}{T}$$

# Adagrad

---

- Divide the learning rate of each parameter by the *root mean square of its previous deviation*

- ◆ Vanilla Gradient descent

- $g^t = \nabla_{\theta} \left( \frac{1}{|S|} \sum_{(x_i, y_i) \in S} \ell(f(x_i; \theta^t), y_i) \right)$
- $\theta^{t+1} = \theta^t - \eta_t g^t$

- ◆ Adagrad

- $r = r + g^t \odot g^t$
- $\theta^{t+1} = \theta^t - \eta_t \frac{1}{\delta + \sqrt{r}} \odot g^t$

# RMSProp

---

## Adagrad

$$r = r + g^t \odot g^t$$
$$\theta^{t+1} = \theta^t - \eta_t \frac{1}{\delta + \sqrt{r}} \odot g^t$$

## RMSProp

$$r = \alpha r + (1 - \alpha) g^t \odot g^t$$
$$\theta^{t+1} = \theta^t - \eta_t \frac{1}{\delta + \sqrt{r}} \odot g^t$$

# Adam

---

## **RMSProp**

$$r = \alpha r + (1 - \alpha) g^t \odot g^t$$

$$\theta^{t+1} = \theta^t - \eta_t \frac{1}{\delta + \sqrt{r}} \odot g^t$$

## **Adam: RMSProp+Momentum**

$$r = \alpha r + (1 - \alpha) g^t \odot g^t$$

$$v = \rho v - \eta_t \frac{1}{\delta + \sqrt{r}} \odot g^t$$

$$\theta^{t+1} = \theta^t + v$$



# Convergence of Adam

---

- Adam may not convergent in some special cases!
- Many provable variants of Adam:
  - ◆ Amsgrad
  - ◆ Adashift
  - ◆ ...

# Adam vs SGD

---

## □ Adam

- ◆ Faster convergence in practice
- ◆ Not sensitive to the learning rate
- ◆ Does not perform well on image classification tasks

## □ SGD

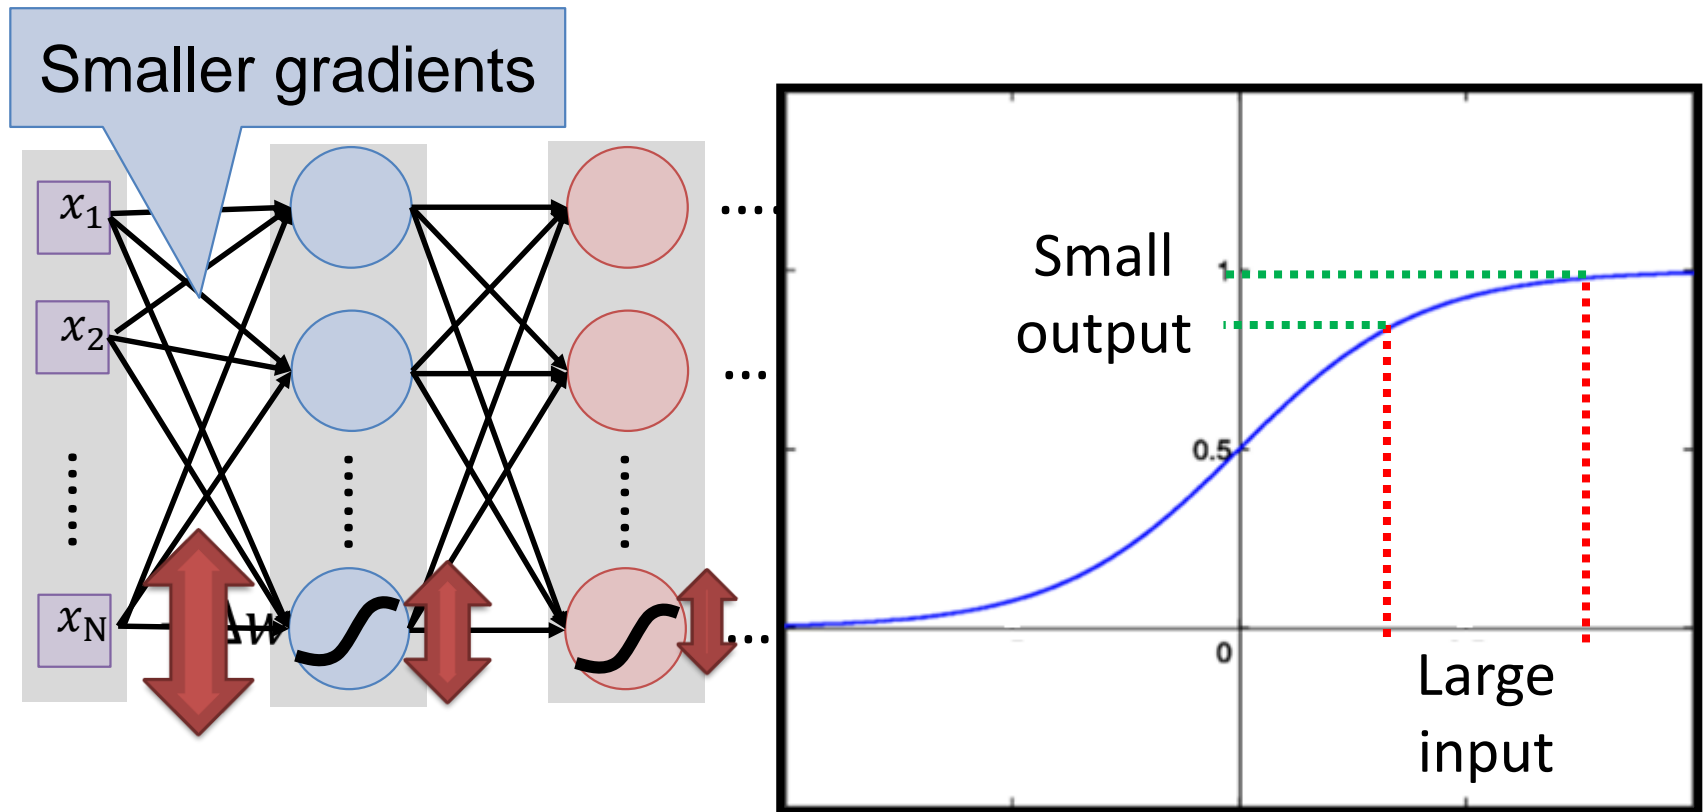
- ◆ Usually slower than Adam
- ◆ Require fine tune of learning rate
- ◆ Has better generalization performance

# Outline

---

- ❑ Model selection
- ❑ Optimization
- ❑ Not work well on training data
- ❑ Not work well on testing data
- ❑ Hyperparameter tuning

# Vanishing Gradient Problem

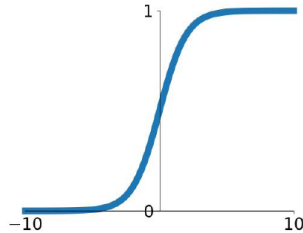


# Activation Functions

---

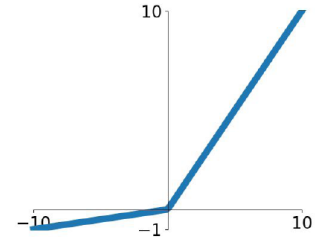
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



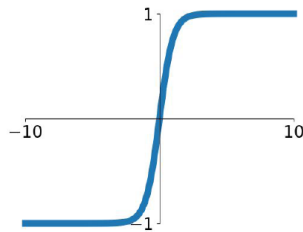
## Leaky ReLU

$$\max(0.1x, x)$$



## tanh

$$\tanh(x)$$

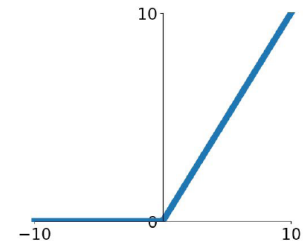


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

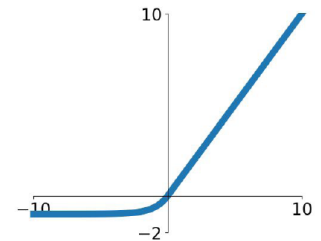
## ReLU

$$\max(0, x)$$



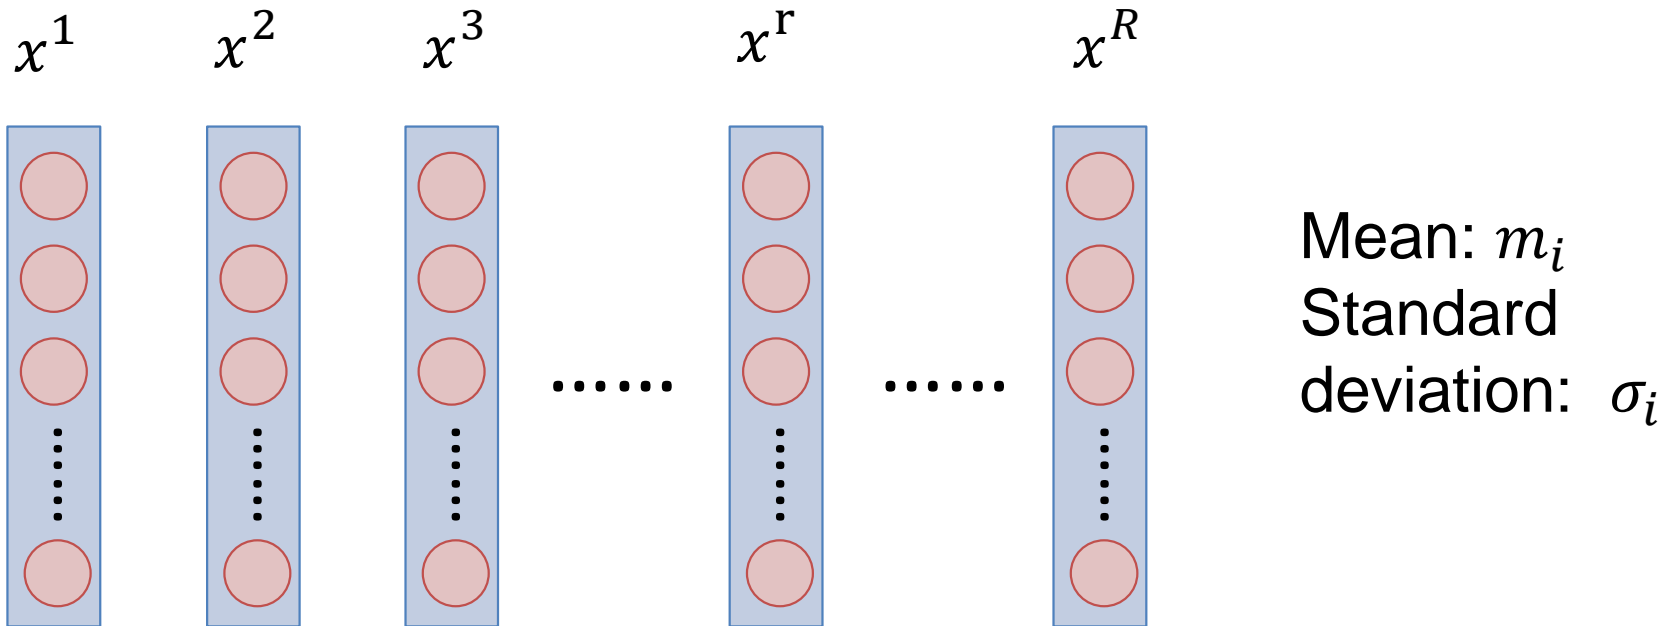
## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Feature Scaling

---



$$x_i^r \leftarrow \frac{x_i^r - m_i}{\sigma_i}$$

# Batch Normalization

---

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

# Outline

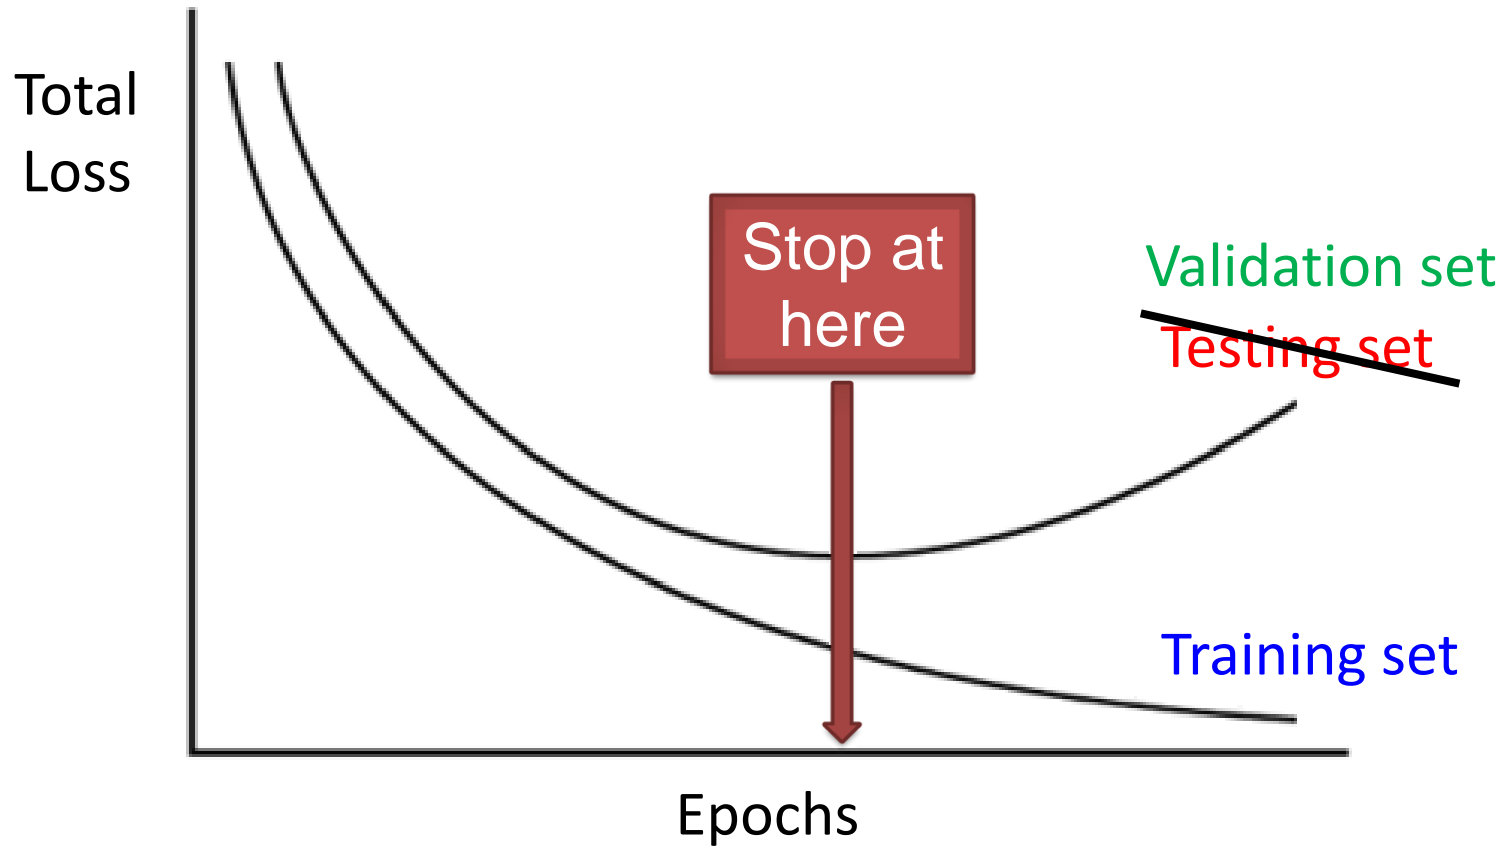
---

- ❑ Model selection
- ❑ Optimization
- ❑ Not work well on training data
- ❑ Not work well on testing data
- ❑ Hyperparameter tuning



# Early stopping

---



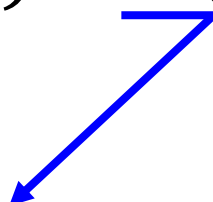
# Regularization

---

- New loss function to be minimized
- Find a set of weight not only minimizing the original cost but also close to zero

$$L'(\theta) = \underline{L(\theta)} + \lambda \underline{R(\theta)} \rightarrow \text{Regularization term}$$

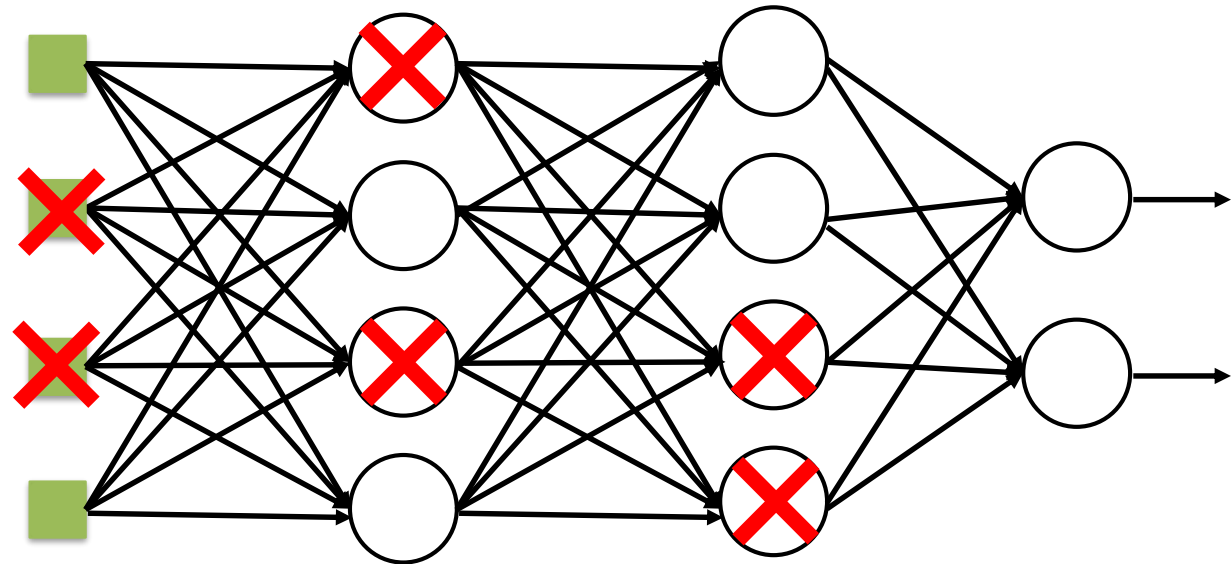
Original loss



# Dropout

---

## Training:

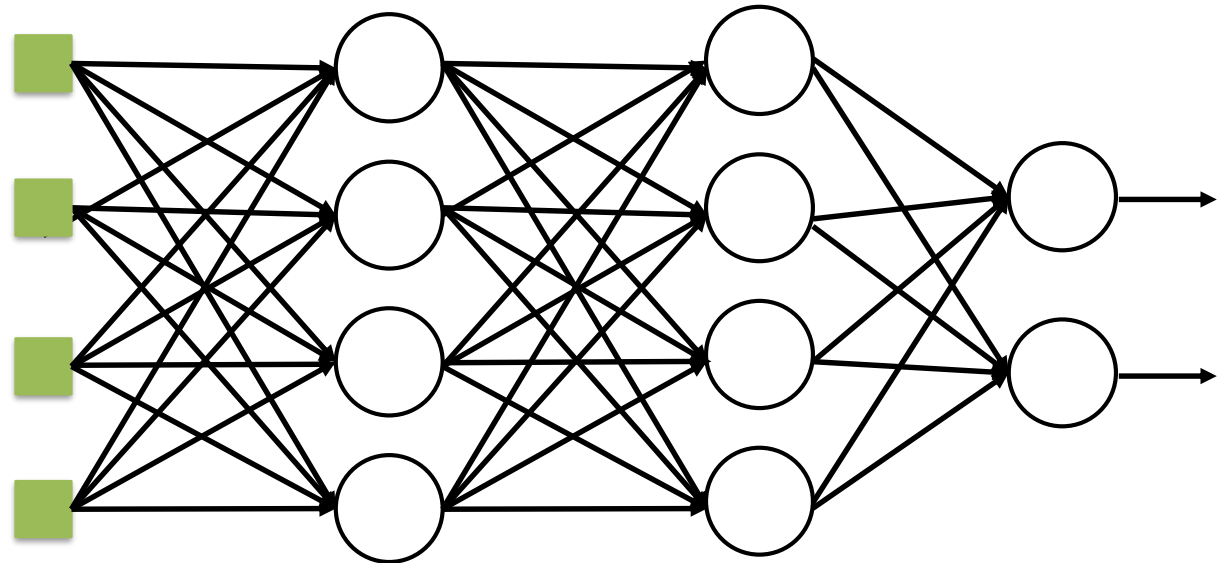


- ◆ Each time before updating the parameters
  - Each neuron has a probability of  $p$  to be dropped

# Dropout

---

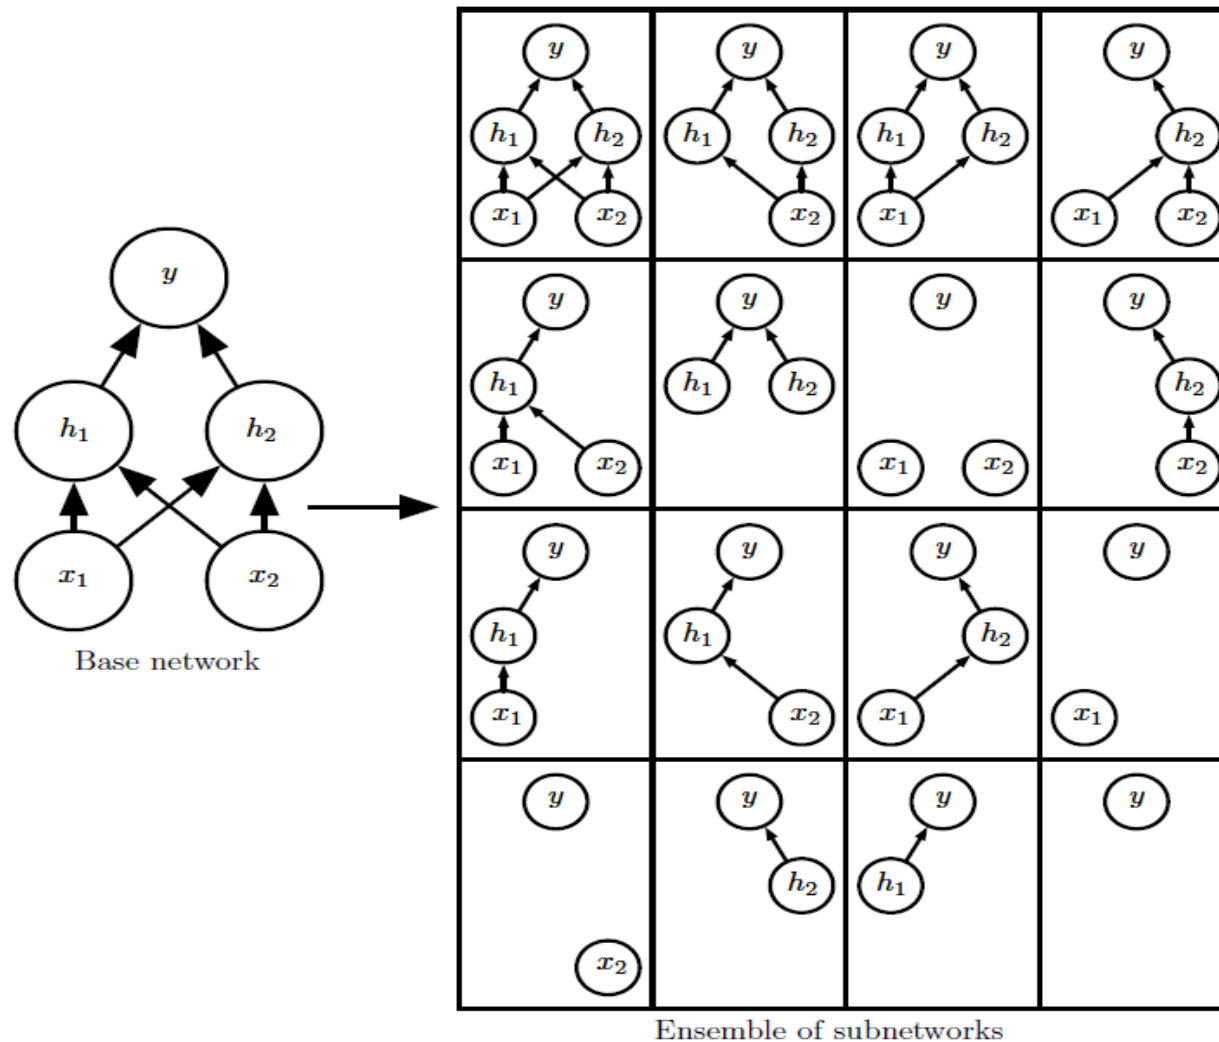
## Testing:



### □ No dropout during testing

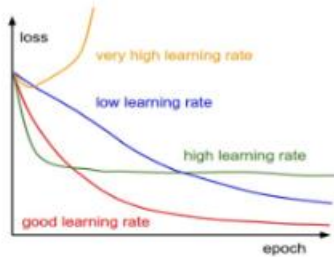
- ◆ If the dropout rate at training is  $p$ , all the weights times  $1 - p$
- ◆ Usually we choose  $p = 0.5$

# Why dropout performs well?



# Hyperparameter tuning

## □ Baby sitting



Loss



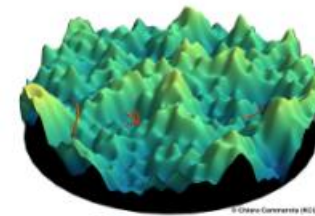
Data



Act/Grad/Filter



Metric



Space

# Grid Search

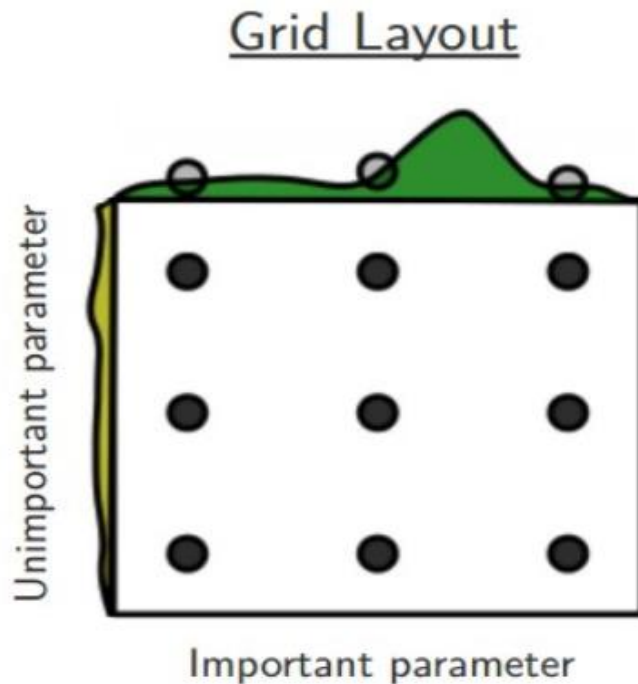
---

## □ Workflow:

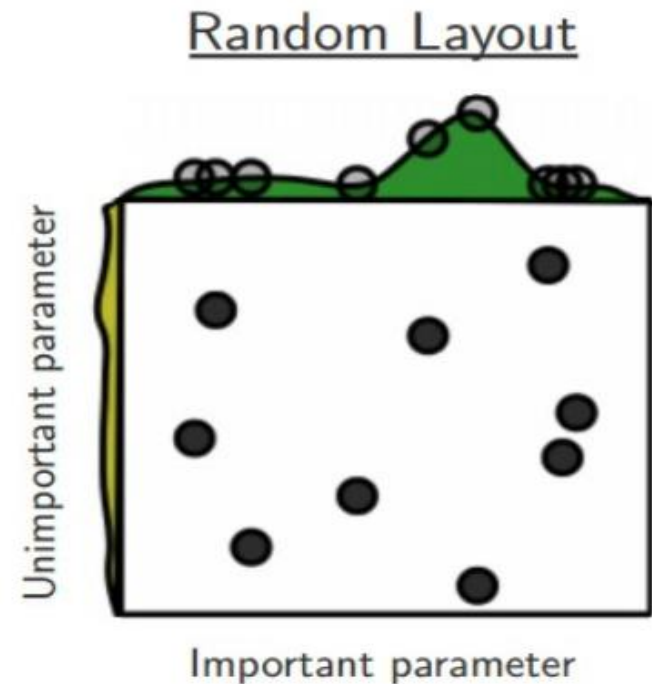
- ◆ Define a grid on  $n$  dimensions, where each of these maps is for an hyperparameter.
  - e.g.  $n = (\text{learning\_rate}, \text{dropout\_rate}, \text{batch\_size})$
- ◆ For each dimension, define the range of possible values:
  - e.g.  $\text{batch\_size} = [8, 16, 32, 64, 128, 256]$
- ◆ Search for all the possible configurations and wait for the results to establish the best one:
  - e.g.  $C1 = (0.1, 0.3, 8) \rightarrow \text{acc} = 92\%$ ,  $C2 = (0.1, 0.35, 8) \rightarrow \text{acc} = 92.3\%$ , etc...

# Grid Search vs Random Search

---



Bad on high dimensional spaces



It doesn't guarantee to find the best hyperparameters

Good on high dimensional spaces

Give better results in less iterations



# Summary

---

- ❑ Supervised Learning: underfitting & overfitting
- ❑ Model selection
- ❑ Optimization
- ❑ Not work well on training data
- ❑ Not work well on testing data
- ❑ Hyperparameter tuning

# Questions

---

